Teradata Bacis

Lesson 06: Teradata Utilities
Multiupload

Module Object

- Introduction about Teradata Utility
- About Multi load
- Supporting Environment
- Multi Load Tasks
- Multi Load Tasks-IMPORT
- Phases of Import Task
- Example of Import Task
- Multi Load Tasks-Delete
- Example of Delete Task
- DELETE Task Differences from IMPORT Task
- Restarting Multiload
- MultiLoad Commands

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 2

Introduction about Teradata Utility

- What is the need of Teradata utilities in Data ware house
  - Quick access to data for more timely decision making.
  - Solutions for the entire spectrum of load requirements from batch to near real time.
  - Unmatched scalability for large volume loads.
  - Fail-proof loads with checkpoint restart capabilities.
  - Proven technology from the data warehouse technology leader.
  - Integration with industry-leading ETL and ELT tools.
- Teradata Utilities
  - BTEQ: Help for Report formatting, Ad hoc query tool, Database administration, Best for small data volumes
  - Multi Load :High-performance data unload in client format. Fast Load: High-performance initial table load.
  - Multi Load: High-performance maintenance operations applies updates to multiple tables in single pass.
  - Apart from these teradata having other utilities like Teradata Parallel Transporter, Tpump e.t.c.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 3

About Multi load

- Multi load is a command driven parallel load utilities , used to load high volume of data to populated or empty
  - teradata tables and views. Teradata MultiLoad executes a series of MultiLoad commands and Teradata SQL statements written in a batch mode job script or interactively entered. The MultiLoad commands provide the session control and data handling specifications for the data transfer operations.
- Features:--
  - Batch mode utility that runs on a server or host system. Supports up to five populated or empty tables.
  - Supports INSERTs, UPDATEs, DELETEs and UPSERTs; typically with batch inputs from a host file. Allows Duplicate rows.
  - Host and LAN support. Full Restart capability.
  - Error reporting via error tables.
  - Support for INMODs MULTILOAD HOST Teradata.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 4

About Multi load

- Features:--
  - Batch mode utility that runs on a server or host system. Supports up to five populated or empty tables.
  - Supports INSERTs, UPDATEs, DELETEs and UPSERTs; typically with batch inputs from a host file. Allows Duplicate rows.
  - Host and LAN support. Full Restart capability.
  - Error reporting via error tables.
  - Support for INMODs MULTILOAD HOST Teradata.

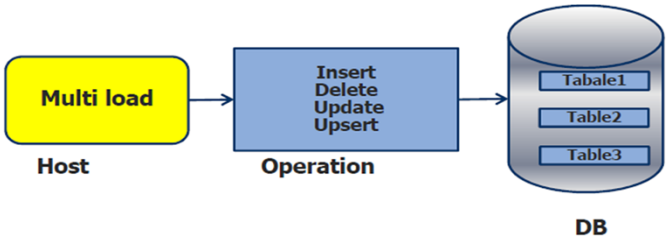

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 5

About Multi Load

- Advantages:-
  - Minimizes the use of the PEs.
  - Gets input data to the AMPs as quickly as possible.
  - Uses multiple-AMP sessions.
  - Uses the parallelism of the AMPs to apply changes.
  - Keeps BYNET activity low with AMP-local processing.
  - Avoids Transient Journaling overhead.
  - Allows Checkpoint/Restartability even with down AMPs.
  - Prevents lengthy rollbacks of aborted jobs.
  - Allows for maximum access to table during processing.
  - Posts errors to special error tables.
  - Provides extensive processing statistics

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 6

Supporting Environment

- The Multi Load utility is supported either on either the mainframe or on network attached system(LAN).
- The LAN environment supports the following Operating Systems:
  - UNIX MP-RAS
  - •Windows 2000
  - •Windows 95
  - •Windows NT
  - •UNIX HP-UX
  - •AIX
  - •Solaris SPARC
  - •Solaris Intel

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 7

Supporting Environment

- The Mainframe (Channel Attached) environment supports the following Operating Systems:
  - MVS
  - VM
- CAUTION: The Teradata RDBMS will only support a maximum of 15 simultaneous Fast Load, Multi Load, or Fast Export utility job

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 8

Multi Load Tasks

- Basically Multiload use the source data and do the operation like Insert, Update, Delete and Up sert and keep in Teradata table.
- There are two types of tasks that Multiload can perform.

1. Import Taslk

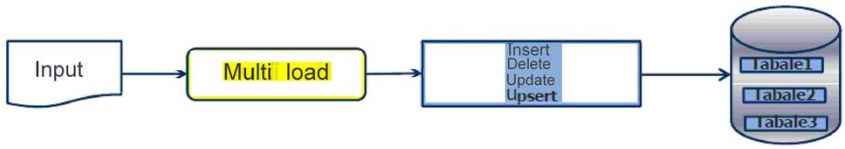

2) Delete Task

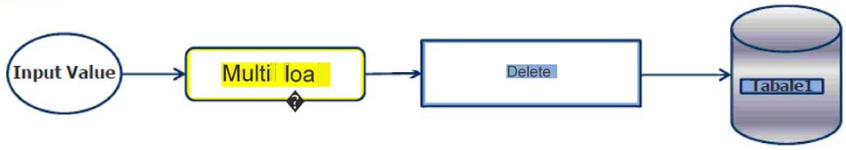

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 9

Multi Load Tasks-IMPORT

- INSERTs, DELETEs, UPDATEs and UPSERTs are allowed.
- Up to a maximum of five tables:–Empty or populated.–NUSIs permitted.
- MultiLoad Import task operations are always primary index operations - however, you are not allowed
- to change the value of a table's primary index.
- Change the value of a column based on its current value.
- Permits non-exclusive access to target tables from other users except during Application Phase.
- Input error limits may be specified as a number or percentage.
- Allows restart and checkpoint during each operating phase.
- IMPORT tasks cannot be done on tables with USI's, Referential Integrity, Join Indexes, Hash
- Indexes, or Triggers.–IMPORT tasks can be done on tables defined with "Soft Referential Integrity".

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 10

Phases of Import Tasks

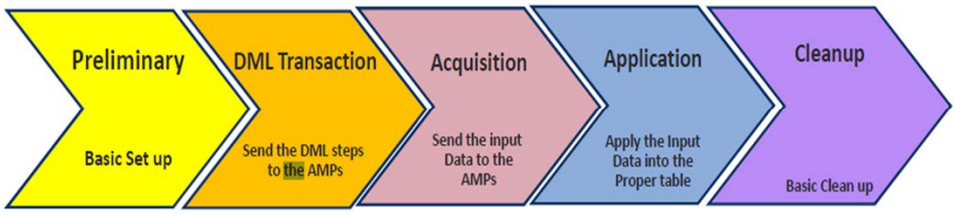

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 11

Multiload: Preliminary Phase

- Preliminary Phase:-
- Validate the Multi load and SQL statements.
  - The first task is to be sure that the SQL syntax and Multi Load commands are valid. Multi Load ensures that it is much better to identify any syntax errors, right up front. All the preliminary steps are automated, so no user intervention is required in this phase.
- Establish all Sessions.
  - Second, all Multi Load sessions with Teradata need to be established. The default is the number of available AMPs.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 12

Multiload: Preliminary Phase

- Preliminary Phase:-
  - Create Work Table Multiload create a work table for its support work. Work Table hold two things: the DML tasks requested and the input data that is ready to APPLY to the AMPs. Default is in user's default database and the work table is named WT_TableName
  - Alternative may be specified as DataBaseName.WorkTableName

Example:

```
BEGIN [IMPORT] MLOAD
TABLES          Employee, Paycheck
 WORKTABLES     util_db.WT_Emp ,
                util_db.WT_Pay
                ...
```

  - Create Error Table Multi Load requires two error tables per target table. The first error table contains constraint violations, while the second error table stores Unique Primary Index violations

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 13

Multiload: Preliminary Phase

- Preliminary Phase:-
  - Error table 1 (ET)
  - Default is the user's database and the table is named ET_Tablename.
    - Contains any errors that occur in the Acquisition Phase.
    - Contains primary index overflow errors that occur in the Application phase
  - Error table 2 (UV)
  - Default is the user's database and the table is named UV_Tablename.
    - Contains Application Phase errors.
    - Uniqueness violations
    - Constraint errors
    - Overflow errors on columns other than primary index

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 14

Multiload: Preliminary Phase

- Preliminary Phase:-
- Example

```
BEGIN [IMPORT] MLOAD
TABLES Employee, PayCheck
WORKTABLES util_db.WT_Emp, util_db.WT_Pay
ERRORTABLES util_db.ET_Emp util_db.UV_Emp,
util_db.ET_Pay util_db.UV_Pay
. . . ;
```

- Apply locks to target tables:
  - The final task of the Preliminary Phase is to apply utility locks to the target tables. Initially, access locks are placed on all target tables, allowing other users to read or write to the table for the time being. However, this lock does prevent the opportunity for a user to request an exclusive lock. Although, these locks will still allow the MultiLoad user to drop the table, no one else may DROP or ALTER a target table while it is locked for loading. This leads us to Phase 2.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 15

Multiload: DML Transaction Phase

- DML Transaction phase:-
  - All of the SQL Data Manipulation Language (DML)
  - statements are sent ahead to Teradata.
  - •Teradata's Parsing Engine (PE) parses the DML and generates a step-by-step plan to
  - •Execute the request. This execution plan is then communicated to each AMP and stored in the appropriate worktable for each target table.
  - Add a USING modifier to the request and host to be filled in from input file.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 16

Multiload: Acquisition Phase

Acquisition Phase:-

- Get the data from host and apply it to appropriate AMP worktables.. → Add Match Tag information to record, Make blocks and send "quick path" to AMPs. Deblock and resend record to "correct" AMP
- Re block and store in worktable off target table → Sort the reblocked records in the work tables, Sort by hash value and sequence to be applied
- Set up transition to the Application phase → Upgrade locks on target tables to Write. Set-table headers for Application phase.

- The lock applied during this phase is:

  - Acquisition lock
    - DML -allows all
    - DDL -allows DROP only

Note: A. Errors that occur in this phase go into the Acquisition Error Table (default name is ET_tablename).
B,. There is no acquisition phase activity for a DELETE Task

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 17

Multiload: Acquisition Phase

- Application Phase:-
  - The purpose of this phase is to write, or APPLY, the specified changes to both the target tables and NUSI sub tables. To accomplish the substitution of data into SQL, when sending the data, the host has already attached some sequence information and five (5) match tags (describe in the next slide) to each data row. Those match tags are used to join the data with the proper SQL statement based on the SQL statement. In addition to associating each row with the correct DML statement, match tags also guarantee that no row will be updated more than once, even when a RESTART occurs.
  - Execute MLOAD for each target table as a single multi-statement request.
    - End of host interaction until end of phase.
    - AMPs independently apply changes to target tables.
    - Executed as a single transaction without rollback.
    - Restart able based on last checkpoint.
    - No transient journal needed.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 18

Multiload: Acquisition Phase

- The lock applied during this phase is:
  - Application lock
  - DML —allows SELECT with ACCESS only
  - DDL —allows DROP only
- Note: Errors that occur in this phase go into the Application Error Table (default name is UV_tablename).

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 19

Multiload: Cleanup Phase

- Cleanup Phase:-
- If the last error code is zero (0), all of the job steps have ended successfully. If this being the case,all error tables made emptied, worktables and the log table are dropped. All locks, both Teradata and MultiLoad, are released. The statistics for the job are generated for output (SYSPRINT) and the system count variables are set. After this, each MultiLoad session is logged off. Below are the summary of the activities:

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 20

Multiload: Cleanup Phase

- All locks are released.
  - Table headers are restored across all AMPs.
  - Dictionary cache of Target Tables is spoiled.
  - Statistics are reported.
  - Final Error Code is reported.
  - Target tables are made available to other users.
  - Work Tables are dropped.
  - Empty Error Tables are dropped.
  - Log Table is dropped (if Error Code = 0).

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 21

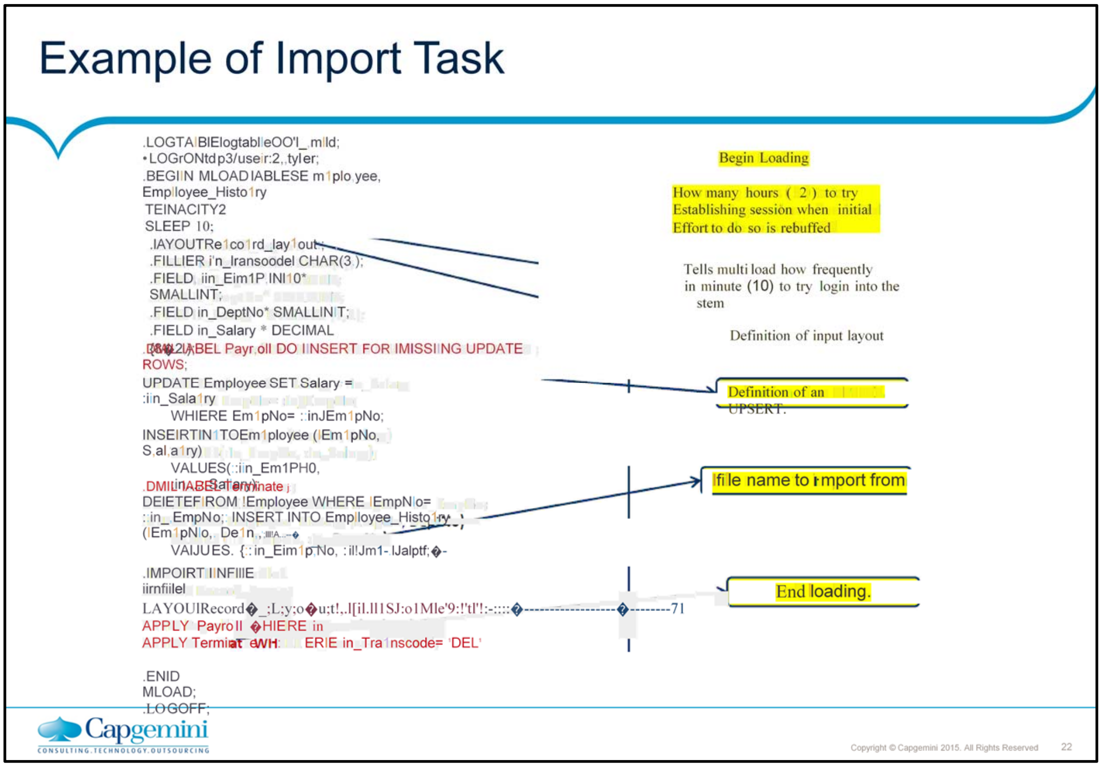
Example of Import Task
.LOGTABLE logtable001_mld;
.LOGON tdp3/user2,tyler;
.BEGIN MLOAD TABLES Employee,
Employee_History
TENACITY 2
SLEEP 10;
.LAYOUT Record_layout;
.FILLER in_Transcode CHAR(3);
.FIELD in_EmpNo
SMALLINT;
.FIELD in_DeptNo SMALLINT;
.FIELD in_Salary DECIMAL
.DML LABEL Payroll DO INSERT FOR MISSING UPDATE
ROWS;
UPDATE Employee SET Salary =
:in_Salary
WHERE EmpNo= :in_EmpNo;
INSERT INTO Employee (EmpNo,
Salary)
VALUES(:in_EmpNo,
:in_Salary);
.DML LABEL Terminate;
DELETE FROM Employee WHERE EmpNo=
:in_EmpNo; INSERT INTO Employee_History
(EmpNo,
VALUES.
.IMPORT INFILE
infile
LAYOUT Record
APPLY Payroll WHERE in
APPLY Terminate WHERE in_Transcode= 'DEL'
.END
MLOAD;
.LOGOFF;
Begin Loading
How many hours (2) to try
Establishing session when initial
Effort to do so is rebuffed
Tells multi load how frequently
in minute (10) to try login into the
stem
Definition of input layout
Definition of an
UPSERT.
file name to import from
End loading.
Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved
22

Multi Load Delete Task

- Not part of Multiload Import, which consider the sql “Delete” statement.
- Not follow the Transient Journal like Import delete.
- No Rollback applies, when job failed.
- Follow the Restartibilty and check points
- Not consider the primary index.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 23

Example of Delete Task

```
..LOGTABLE ILogta bl1e003_mid
.LOGON tdlp3/user2 ,tyler
,.  B.EGI N DELETE MLOAD TA.BLES Employee
.LA.VOUT R1emove;
.FIELD inTermdlate* INITEGER;
  DELETE FROM Employ ee WHERE le rm_date >
:in_J�ermdate;
.IMPORT INIFII-LEinfUe2LAYOUT Remov e;
.END MI.JOA.D;
.LOGOFIF;
```

Begin stmt for delete records

Delete stmt

End Mload

Capgemini CONSULTING.TECHNOLOGY.OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved

DELETE Task Differences from IMPORT Task

- DELETE tasks operate very similarly to IMPORT tasks with some differences:
  - Deleting based on a equality UPI Value is not permitted.
  - A DML DELETE statement is send to each AMP with match tag parcel.
  - No Acquisition phase.
  - Application Phase reads all target blocks and deletes qualifying rows.
- Multiload Delete vs SQL Delete.
  - Multi load Delete is Faster as compared to SQL Delete.
  - Multi load Delete is restartable.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 25

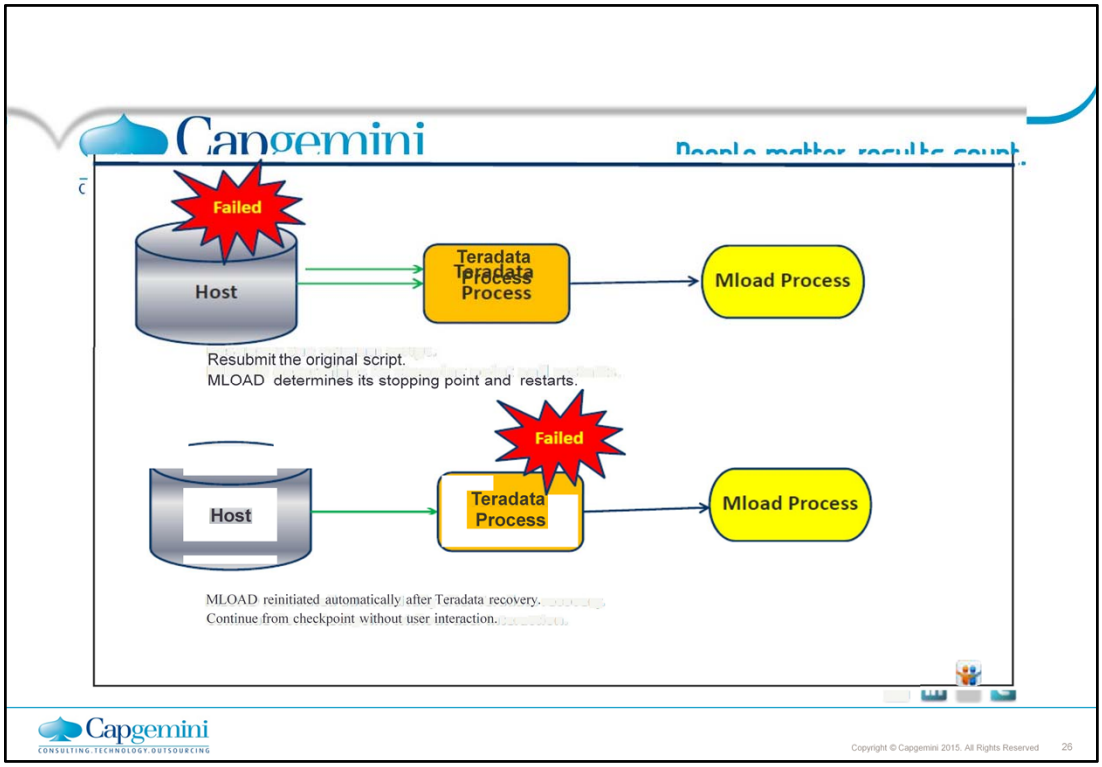
Capgemini
Failed
Host
Teradata Process
Mload Process
Resubmit the original script.
MLOAD determines its stopping point and restarts.
Failed
Host
Teradata Process
Mload Process
MLOAD reinitiated automatically after Teradata recovery.
Continue from checkpoint without user interaction.
Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved
26

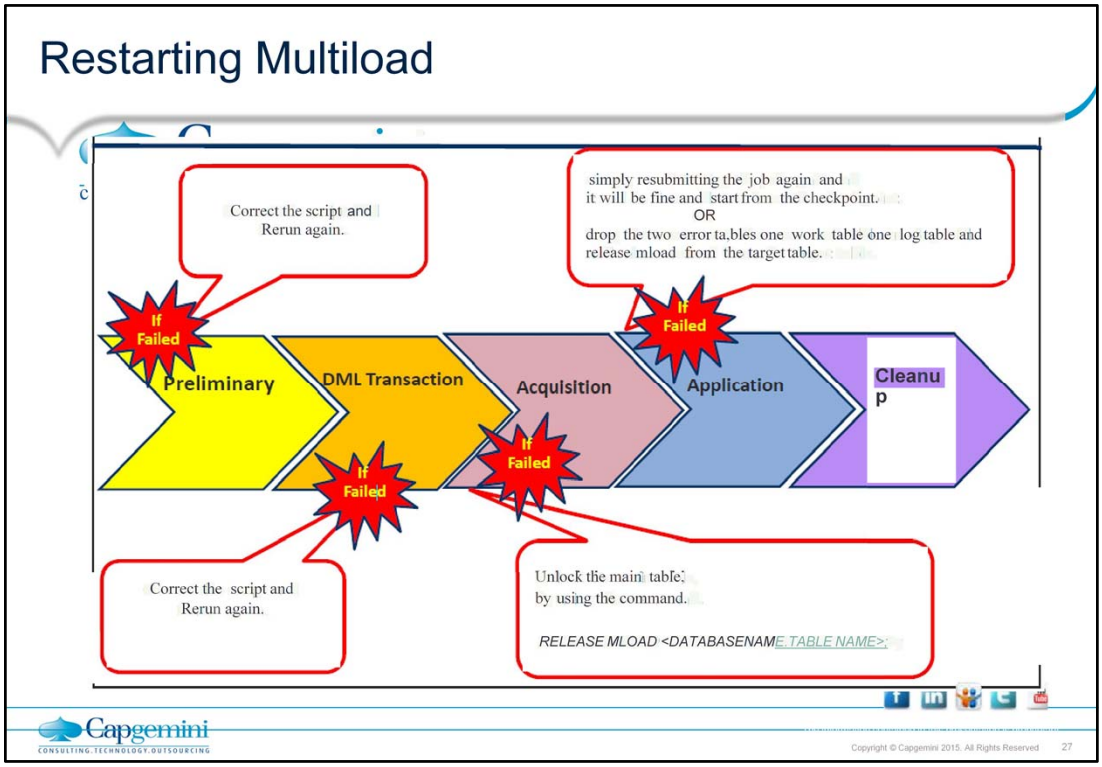
Restarting Multiload
Correct the script and Rerun again.
simply resubmitting the job again and it will be fine and start from the checkpoint.
OR
drop the two error ta,bles one work table one log table and release mload from the target table.
If Failed
Preliminary
DML Transaction
Acquisition
Application
Cleanup
Correct the script and Rerun again.
Unlock the main table; by using the command.
RELEASE MLOAD <DATABASENAME.TABLE NAME>;
Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING
Copyright © Capgemini 2015. All Rights Reserved
27

MULTILOAD COMMAND

.BEGIN =IMPORT] MLOAD .BEGIN DE1LETE MLOAD	This command communicates directmy with Te!J'"adata to specify if the MulbiLo.ad mode is going to be IMPORT or DELETE. Note that the Word £MPORT is. opt:ional iin ilthe synl:ax because it. iis. the DIE FAULT" but DELETE is required. *!Ve recommend usin9* Ehe. *word IMPORT* to make tine coding consistent and easier for others t:o read. Any parameters for the load, such as error limits or checkpoints will be included under the . BEGIN ccnnrnerrd, too. It is importa111t rO know which commands or parameters are optional since, if you do not include them, Mult1iLoad may supply defaults that may impact your load.
.DML LABEL	The DML IL.ABEL defines treatment options and labels for the application {APPLY) of data for tihe INSERT, UPDATE, UPSERT and DELETE op ere tioris. A LABEIL is simply a name for a requested SQL activity. The LABEL is defined first, and th12n referencsd later n the APPLY clause. This Mu!!!:iLoad to f!:n!sh the I\PDLY cper-.a tiorrs ·.dth the ch.ariges to t1,1e desiqnated databases and tables.
.FIELD	This defines *a* column of the data sourrce record that will be sent to the Teradata database via SQL. \.i\llhen writing the script yo._1 must rnciude *a* FIELD for each data field you need in SQL. This c ornrriarrd is used with the L...,.YOUT command.
.FILLER	Do not assume that MultrLoad ha1s somehow uncovered rrurcb of vfi·hat y.o._, used in your terr1n papell"S at the university! FILLER defines a field that is accounted for as part of the data source's row format,, but: is *not:* sent to tlle Teradata DBS. It is used with the LAYOUT comrnarid:
.LAYOUT	LAYOUT defiines the format of the INPUT DATA record so Teradata kno rr1s what to expect. If one record 1s not: large enough,, you can concat:.enate multiple by using the LAYOUT parameter CONTINUEIF t:o tell which value to perform for the conc.atenation. Another option is INDICATORS, ri1hich is used to represent: nulls by using the bitn,ap (1 bit per field) at. the front of tlle data record.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 28

MULTILOAD COMMAND

.LOGON	This specifies the username or LOGON string that will establish sessions for Multiload with Teradata.
LOGTABLE	This support command names the name of the Restart Log that will be used for storing CHECKPOINT data pertaining to a load. The LOGTABLE is then used to tell Multiload where to RESTART1 should that be necessary. It is recommended that this command be placed *before* the .LOGON command.
.LOGOFF	This command terminates any sessions established by the LOGON command
.IMPORT	This command defines the INFILE, the LAYOUT to use and where to APPLY the data to SQL.
.SET	Optionally, you can SET utility variables. An example would be {.SET DBName TO CD_ILLTest}.
.SYSTEM	This interrupts the operation of Multiload in order to issue commands to the local operating system.
.TABLE	This is a command that may be used with the .LAYOUT command. It identifies a table whose columns (name, their order and data types) are to be used as the field names and data descriptions of the data source records.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 29

MULTILOAD Parameters

Parameter	Description
SESSIONS	This refers to the number of SESSIONS that should be established with Teradata. For Multiload, the optimal number of sessions is the number of AMPs in the system, plus two more. (One for Control functions arid other is for the backup purpose). You can also use MAX or MIN, which automatically use the maximum or minimum number of sessions to complete the job. If you specify nothing, it will default to MAX.
SLEEP	Tells MultiLoad how frequently, in minutes, to try logging on to the system.
TENACITY	Tells MultiLoad how many hours to try logging on when its initial effort to do so is rebuffed.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 30

Limitation of MultiLoad

- Though MultiLoad is a very powerful utility; it has following limitations:
  - 1. MultiLoad Utility doesn't support SELECT statement.
  - 2. Concatenation of multiple input data files is not allowed.
  - 3. MultiLoad doesn't support Arithmetic Functions i.e. ABS, LOG etc. in Mload Script.
  - 4. MultiLoad doesn't support Exponentiation and Aggregator Operators i.e. AVG, SUM etc. in Mload Script.
  - MultiLoad doesn't support USIs (Unique Secondary Indexes), Referential Integrity,
  - Join Indexes, Hash Indexes and Triggers.
  - Import task require use of PI (Primary Index).

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 31

MultiLoad Review Questions

- Answer True or False.
  - 1. With MultiLoad, you can import data from the host into populated tables.
  - 2. MultiLoad cannot process tables with USIs or Referential Integrity defined.
  - 3. MultiLoad allows changes to the value of a table's primary index.
  - 4. MultiLoad allows you to change the value of a column based on its current value.
  - 5. MultiLoad permits non-exclusive access to target tables from other users except during Application Phase

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 32

MultiLoad Review Questions

- Match the MultiLoad Phase in the first column to its corresponding task in the second column.

1._____Preliminary	A. Acquires or creates Restart Log Table.
2._____DML Transaction	B. Locks are released.
3._____Acquisition	C. Applies (loads) data to the work tables.
4._____Application	D. Execute mload for each target table as a single multi-statement request.
5._____Cleanup	Stores DML steps in work tables.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 33

Answers of MultiLoad Review Questions

- Match the MultiLoad Phase in the first column to its corresponding task in the second column.

1.___A___Preliminary	A. Acquires or creates Restart Log Table.
2.___E____DML Transaction	B. Locks are released.
3.___C___Acquisition	C. Applies (loads) data to the work tables.
4.___D___Application	D. Execute mload for each target table as a single multi-statement request.
5.___B___Cleanup	Stores DML steps in work tables.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 34

ANSWERS OF MULTILOAD REVIEW QUESTION

- With MultiLoad, you can import data from the host into populated tables.
  - True
- MultiLoad cannot process tables with USIs or Referential Integrity defined.
  - True
- MultiLoad allows changes to the value of a table's primary index
  - Flase
- MultiLoad allows you to change the value of a column based on its current value.
  - True
- MultiLoad permits non-exclusive access to target tables from other users except during Application Phase.
  - True

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 35

MultiLoad LAB Exercise

- Purpose:
  - In this lab, you will use MultiLoad utility to insert the data rows in OLAP_EXAMPLE table. We use the Linux/Unix environment to execute the script.

- What you need:
  - The required table and data file.
- Tasks:
  - 1.Prepare a MultiLoad script which inserts rows into the table using the redefinition feature of MultiLoad.

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 36

MultiLoad LAB Exercise

- Tasks:
  - 1.Prepare a MultiLoad script which inserts rows into the table using the redefinition feature of MultiLoad.
  - 2. Ensure that the table has been created in the database properly.
  - 3. Run the script.

- Note: The Trans_Date is exported with an ANSI Date Format of 'YYYY-MM-DD' and a data type of CHAR(10).

Capgemini
CONSULTING.TECHNOLOGY.OUTSOURCING

Copyright © Capgemini 2015. All Rights Reserved 37